


Exercise 3.4 — Run scikit-learn LogisticRegression

The scikit-learn package provides a [LogisticRegression](#) object to perform logistic regression.

Write code to fit a **LogisticRegression** model using the same training matrix X that you defined as part of Exercise 3.3. There are only two steps:

1. Create the `LogisticRegression` object. Do not fit an "intercept" and do not include any regularization penalty (see [documentation](#)).
2. Fit the `LogisticRegression` object to the training matrix X and targets y . Use the object's `fit` method.

The variable holding a reference to your `LogisticRegression` object should be called `logistic_model`, so that your answer can be checked.

A tweet regarding the fact that scikit-learn's `LogisticRegression` object applies regularization (a weight penalty) "by default":


In [334]:

```
# Your code here. Aim for 2 lines.  
logistic_model = sklearn.linear_model.LogisticRegression(penalty='none', fit_intercept=False)  
logistic_model.fit(X, y)
```

Check your answer by running the code cell below.

In [335]:

```
assert 'logistic_model' in globals(), "You didn't create a variable named 'logistic_model'!"  
assert isinstance(logistic_model, sklearn.linear_model.LogisticRegression), "Expected a LogisticRegression instance"  
assert hasattr(logistic_model, 'coef'), "No model coefficients yet! You didn't fit the model to any data!"  
assert logistic_model.intercept_ == 0.0, "You forgot to disable fitting of the intercept!"  
assert np.allclose(logistic_model.coef_, [[18.5291137, 10.49283446]]), "The parameters seem incorrect! Not 1-Bit!"  
print("Correct!")
```

Notice that the model parameters (coefficients) found by the `LogisticRegression` are much larger than those found by your gradient descent solver. That is only because scikit-learn uses a more powerful optimization algorithm and can learn very sharp decision boundaries in fewer steps than mere gradient descent can. If you increase your `num_steps` argument your solver will find similarly large coefficients.

Plot several **LogisticRegression** predictions at once by running the code cell below.

In [336]:

```
x0 = np.ones(50) # A column of 1s so that the bias term w[0] gets added  
x1 = np.linspace(-5, 5, 50) # A column of x values ranging from [-5, 5]  
X_test = np.column_stack([x0, x1]) # A 20x2 matrix where X[i,:], is the ith x vector  
y_test = logistic_model.predict_proba(X_test) # Evaluate all x values and get two probabilities back (class 0, class 1)  
plt.scatter(x1, y_test[:,1], 10, 'r') # Plot probability of class 1 only  
plt.xlabel("x1")  
plt.ylabel("y1")  
plt.title("Sample predictions for LogisticRegression model")
```



In []: