

SOEN 363 - Database Systems for Software  
Engineers  
**Assignment 3**


Ali Hanni - 40157164

*Gina Cody School of Computer Science and Software Engineering  
Concordia University, Montreal, QC, Canada*

Winter 2022

## Question 1

- (a) We assume that we have an ISAM structure with 4 entries per page. While 25 is directly inserted in a page, 35 will be an overflow node. New ISAM structure is shown bellow.



- (b) The correct tree is Tree B  
 (c) The correct tree is Tree A  
 (d) The correct tree is Tree A

## Question 2

- (a) First, we try to find the order of the  $B^+$  tree  $d$ . We know that the maximum number of keys is  $2d$  and the maximum number of pointers to children is  $2d + 1$ . Knowing that key size is 50-byte, that pointer size is 8-byte, and that page size is 2000 bytes, we have to maximize the value of  $d$  such that:

$$\begin{aligned}(2d) * 50 + (2d + 1) * 8 &\leq 2000 \\ 100d + 16d + 8 &\leq 2000 \\ 116d &\leq 1992 \\ d &\leq 17.17 \\ d &= 17\end{aligned}$$

Hence, when maximizing the value of  $d$  we obtain  $d = 17$ . We know that we can have 34 keys and 35 pointers. Then, we know that a record on any given page holds a key field (50-byte) and a pointer (8-byte). Since the size of one page on the disk is 2000 bytes, we can find the maximum number of entries in a page by doing:

$$\lfloor 2000 / (50 + 8) \rfloor \lfloor 2000 / 58 \rfloor \lfloor 34.48 \rfloor = 34$$

Hence, there is a maximum of 34 entries per page. We can proceed and compute the number of levels in the tree by doing:

$$\begin{aligned}\lceil \log_{35}(20000/34) + 1 \rceil \\ \lceil \log_{35}(588.2353) + 1 \rceil \\ \lceil 1.79367 + 1 \rceil \\ \lceil 2.79367 \rceil = 3\end{aligned}$$

There are 3 levels in the  $B^+$  tree described.

- (b) Since we deal with a dense  $B^+$  tree, we know that the nodes are filled at each level we know that there are  $\lceil 20000/34 \rceil = 589$  leafs on level 3. Since each full node has  $2d + 1 = 2 * 17 + 1 = 35$  children, there are  $\lceil 589/35 \rceil = 17$  nodes on level 2 and a single node on level 1 which is the root of the tree.
- (c) We try to find the order of  $B^+$  but this time with a key size of 10-bytes and a page size of  $2000 * 0.70 = 1400$ :

$$\begin{aligned}(2d) * 10 + (2d + 1) * 8 &\leq 1400 \\ 20d + 16d + 8 &\leq 1400 \\ 36d &\leq 1392 \\ d &\leq 38.66 \\ d &= 38\end{aligned}$$

We can have 76 keys and 77 pointers. Since the size of one page on the disk is 2000 bytes, but pages are only 70% full, we can find the maximum number of entries in a page by doing:

$$\lfloor 1400/(10 + 8) \rfloor \lfloor 1400/18 \rfloor \lfloor 77.7778 \rfloor = 77$$


Hence, there is a maximum of 34 entries per page. We can proceed and compute the number of levels in the tree by doing:

$$\begin{aligned} & \lceil \log_{77}(20000/77) + 1 \rceil \\ & \lceil \log_{77}(259.74026) + 1 \rceil \\ & \lceil 1.2799 + 1 \rceil \\ & \lceil 2.2799 \rceil = 3 \end{aligned}$$

There are 3 levels in the  $B^+$  tree described.

### Question 3

(a) Here is the  $B^+$ -tree index :



The position of the tuples in the file (page #, slot #) are used to identify the different tuples. Note that for practical reasons, when a certain key (GPA) corresponds to a number of distinct entries that is too large, the list is replaced with a variable. List variables definitions are bellow:

$L1 = [(1, 3), (1, 4), (2, 1), (2, 2), (2, 3), (2, 4), (3, 1)]$

$L2 = [(1, 2), (3, 2), (3, 4), (4, 2)]$

(b) 1. If the tuples in  $f$  are sorted, they will appear in the following order:

<1.8 (1, 1)>

<2.2, (1, 2)>

<3.2, (1, 3)>

<3.4 (1, 4), (2, 1), (2, 2), (2, 3), (2, 4), (3, 1), (3, 2)>

<3.5 (3, 3)>

<3.8, (3, 4), (4, 1), (4, 2), (4, 3)>

Hence, with the use of a  $B^+$ -tree index on GPA, the system would need to access 3 pages (pages 1, 2, and 3) before returning all values of tuples with a GPA between 3.0 and 3.5 inclusive.

2. If the tuples in  $f$  are not sorted, they will appear in the following order:

<1.8 (1, 1)>

<2.2, (4, 3)>

<3.2, (3, 3)>

<3.4 (1, 3), (1, 4), (2, 1), (2, 2), (2, 3), (2, 4), (3, 1)>

<3.5 (4, 1)>


<3.8, (1, 2), (3, 2), (3, 4), (4, 2)>

Hence, the system would need to access all 4 pages before returning all values of tuples with a GPA between 3.0 and 3.5 inclusive.

3. We realize that the number of pages accessed is 33% higher when tuples are unsorted in comparison to when they are sorted. This percentage quickly increases as the distribution of GPA widens and the number of tuples and tuples per page increases.

## Question 4

(a) The table resulting from the insertion of above records is shown below:



- (b) The global depth is 3.
- (c) There will be 4 buckets.
- (d) i, h, and b. The bucket has a local depth of 3.
- (e) g, e, and c. The bucket has a local depth of 1.
- (f) **a**). The number of bits to be used in the hash function are kept in the global depth.
- (g) No.
- (h) If the local depth of a bucket is equal to the global depth of the directory, then the bucket is pointed to by exactly one directory entry.